

CSE446: Blockchain & Cryptocurrencies

Lecture - 8: Bitcoin-2



Inspiring Excellence

Agenda

- Bitcoin components
 - Node & Network
 - Bitcoin Blockchain
 - Transaction
 - Block
 - Blockchain

Bitcoin P2P network

- Bitcoin nodes communicate in a decentralised fashion, meaning that no single entity or node is superior, all nodes are equal
 - Ad-hoc protocol (runs on TCP port 8333)
 - Ad-hoc network with random topology
- New nodes can join at any time
- Forget non-responding nodes after 3 hr

Bitcoin P2P network

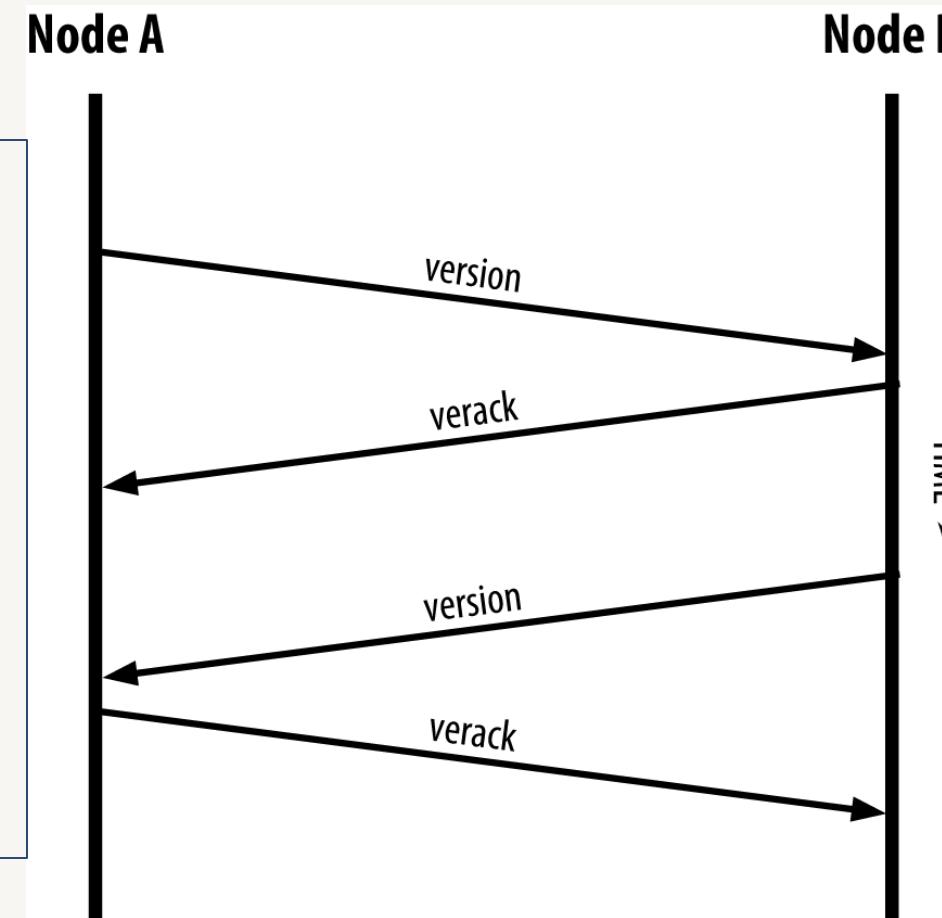
- To communicate, they need to have clear rules
 - How to find other nodes (bootstrapping)
 - How to send and receive transactions
 - How to send and receive blocks
 - How to sync the blockchain
- The basic network uses a peer-to-peer gossip protocol for
 - Node discovery, node status maintenance
 - Messages about new blocks or transactions

Node discovery/bootstrapping

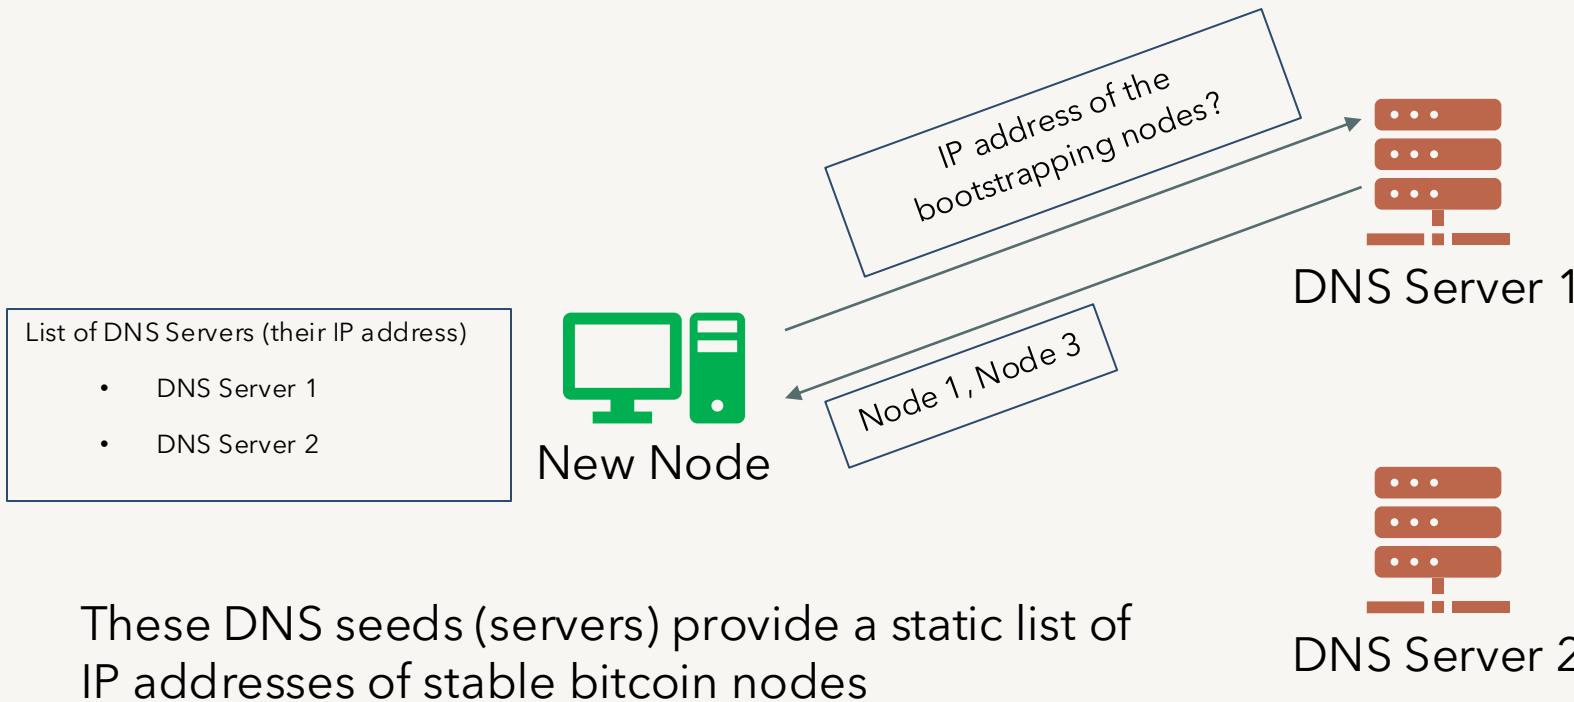
- Adding a new node into the network is called bootstrapping
- The new node needs to discover other nodes in the network to connect to the P2P network
- How does it know who to connect to?
 - Hard-coded DNS-services which offer IP-addresses of nodes
 - Hard-coded seed addresses (last resort)
 - Command-line provided addresses
 - Text-file provided addresses

Node discovery/bootstrapping

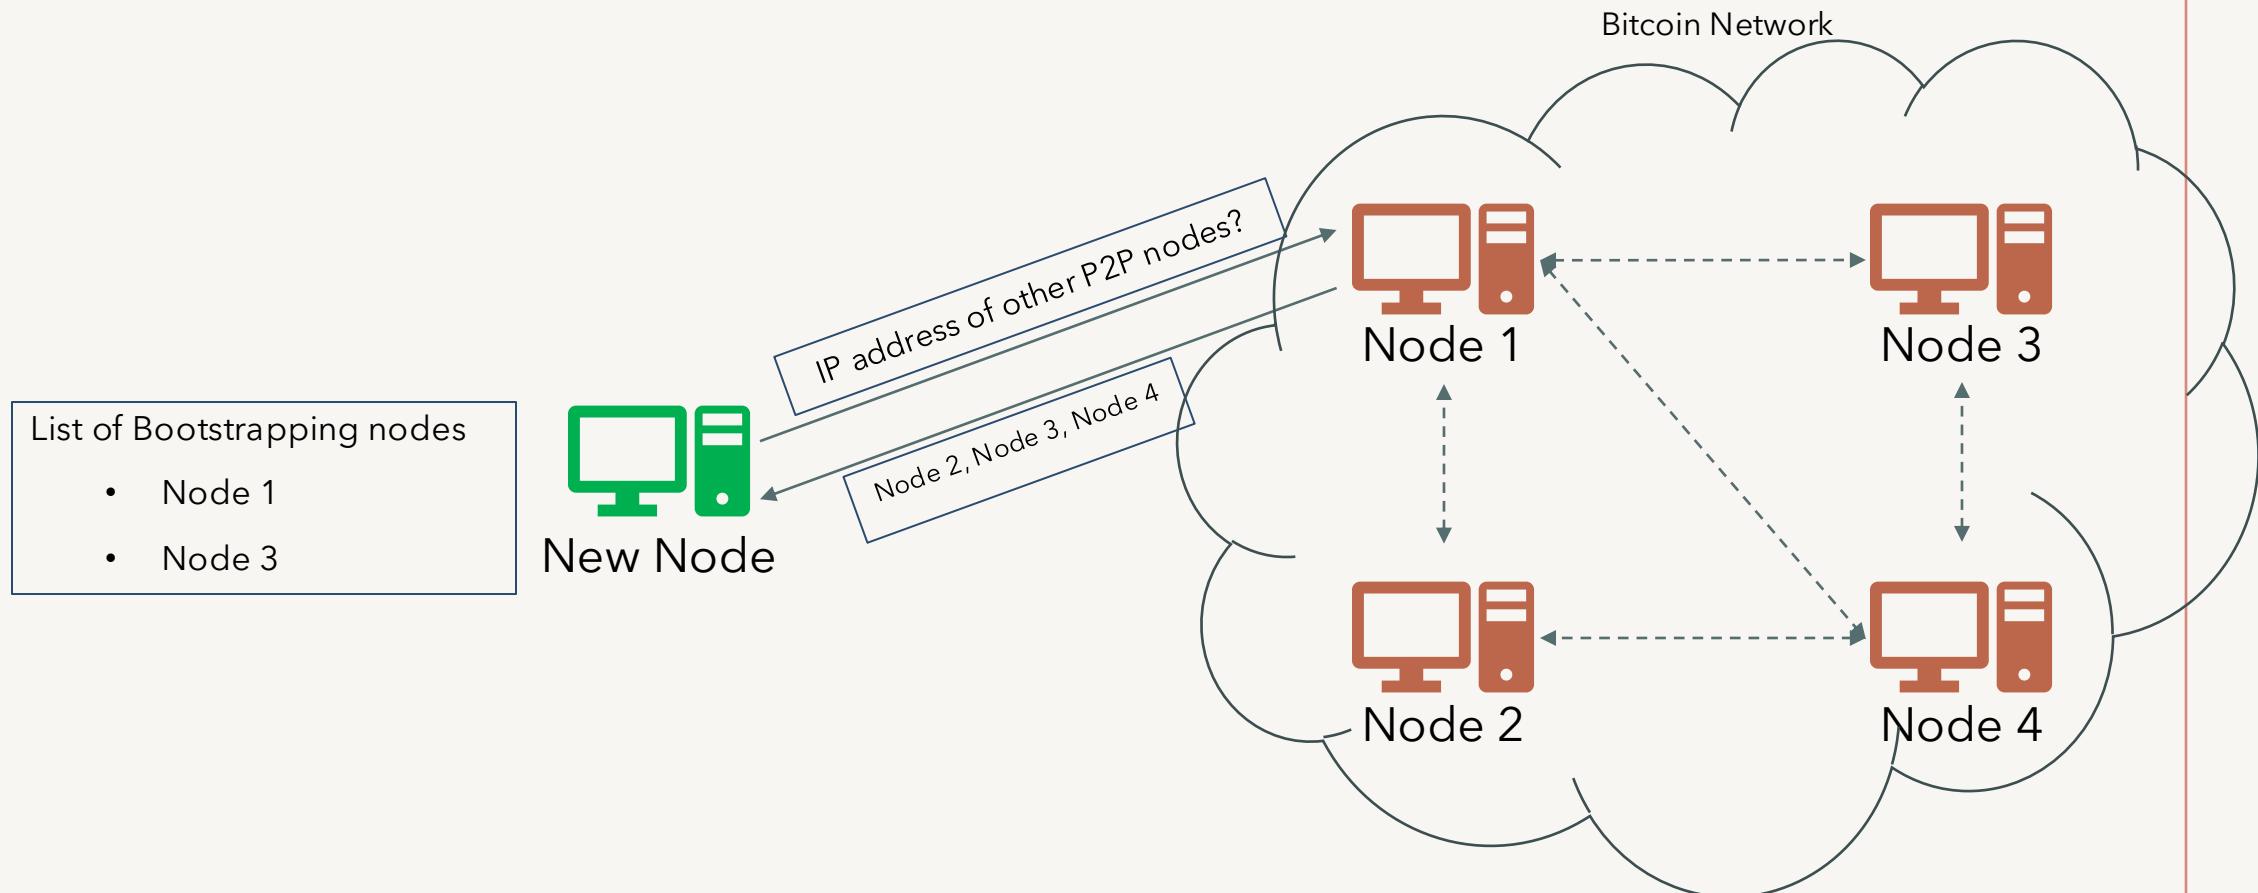
- The first message to another Bitcoin peer is the version message
- Using version each node check if the other node is compatible
- If compatible, the other node sends the version acknowledgement (verack) message



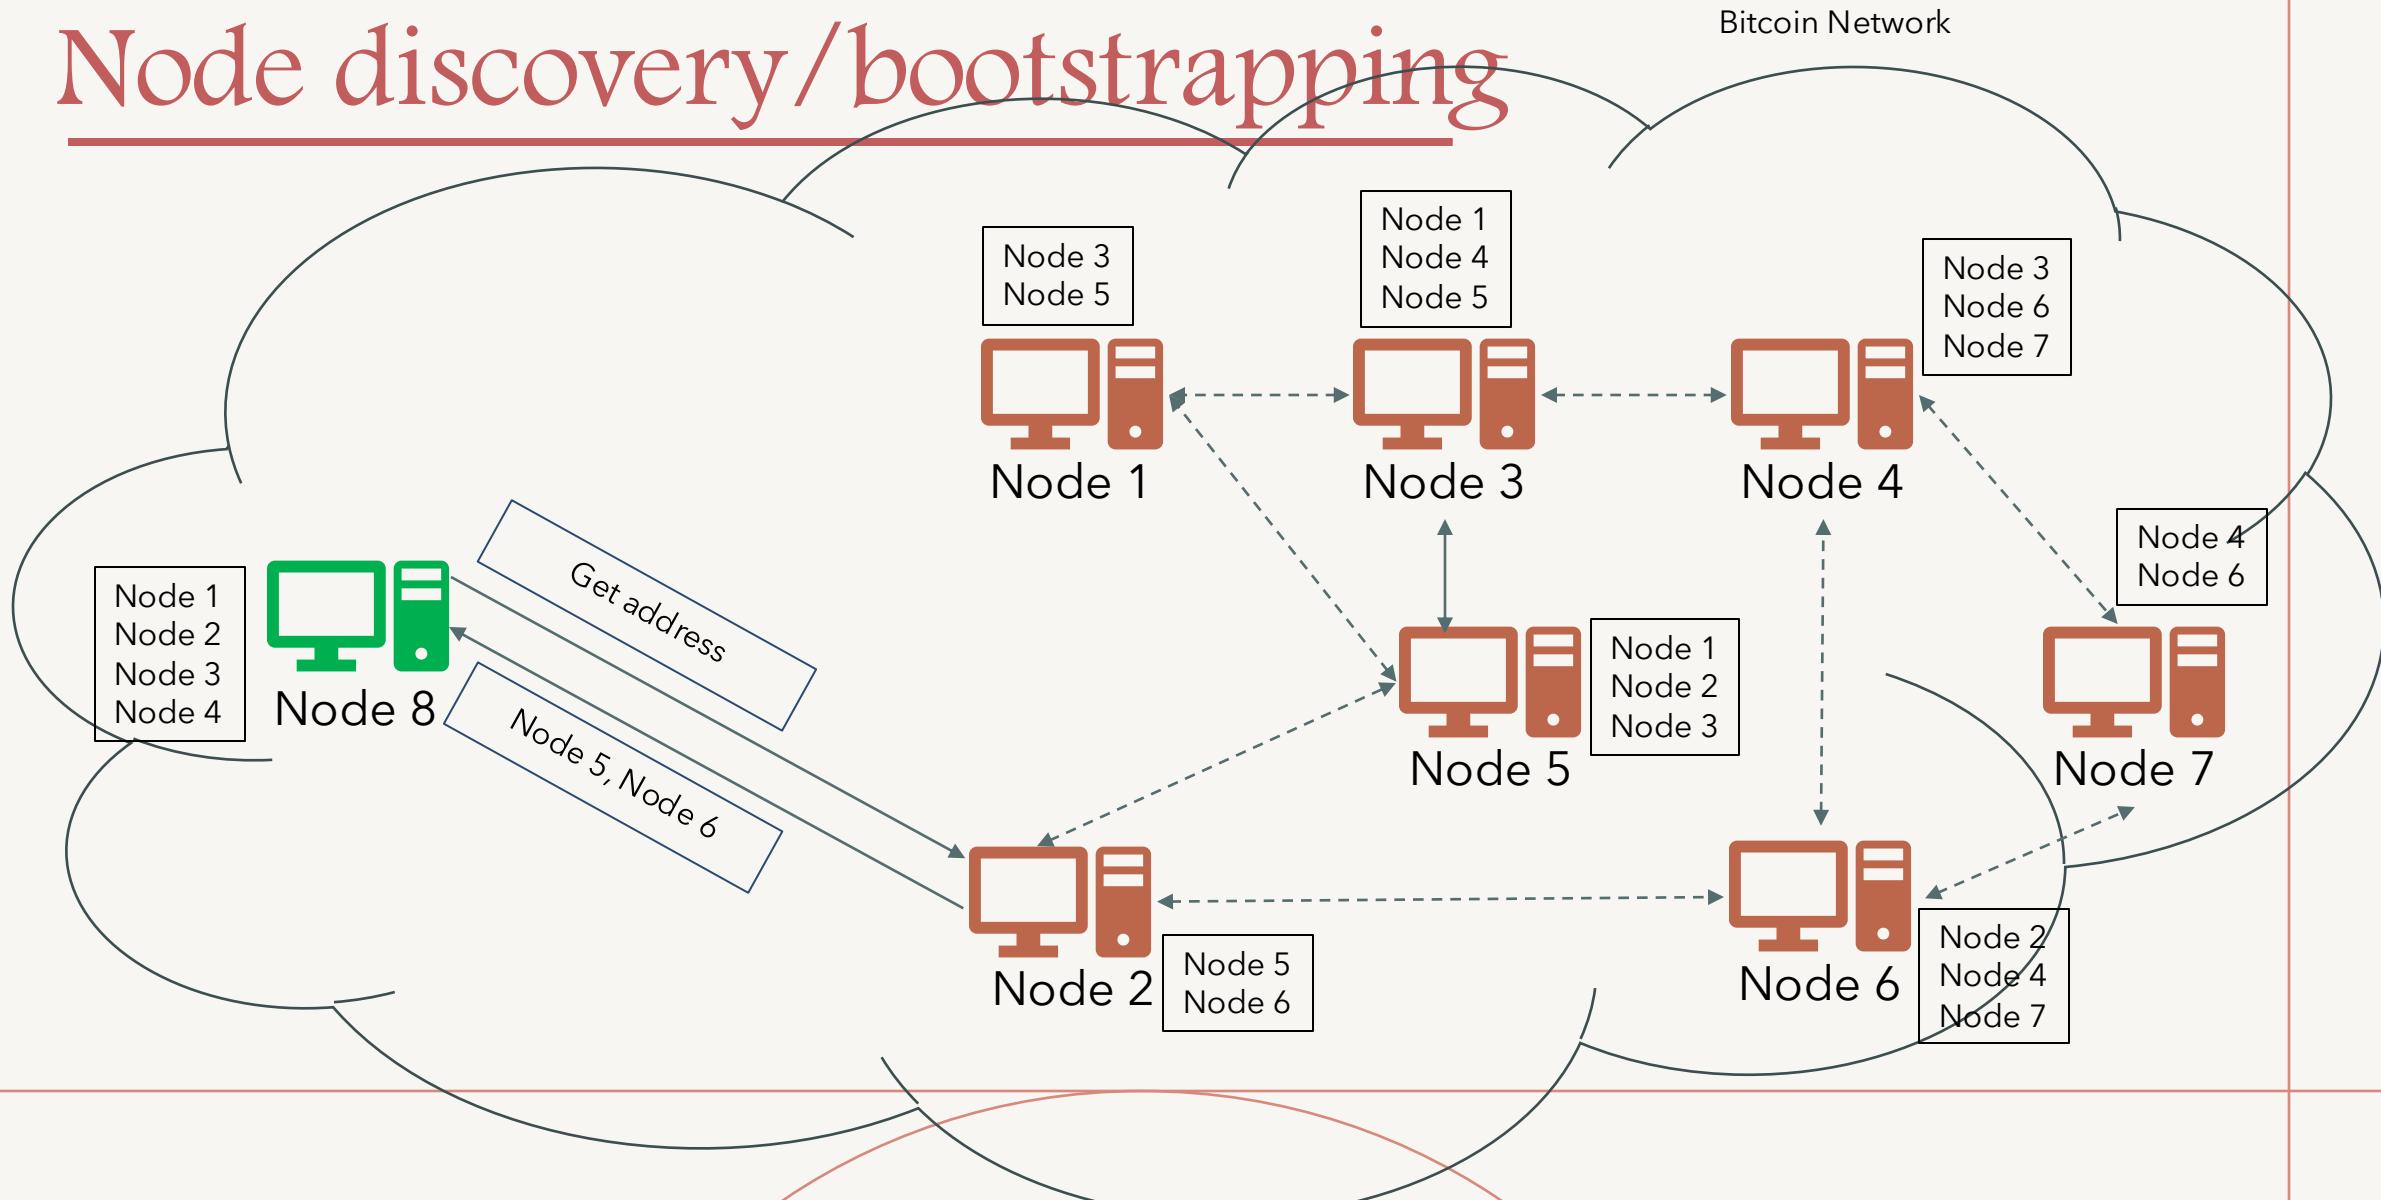
Node discovery/bootstrapping



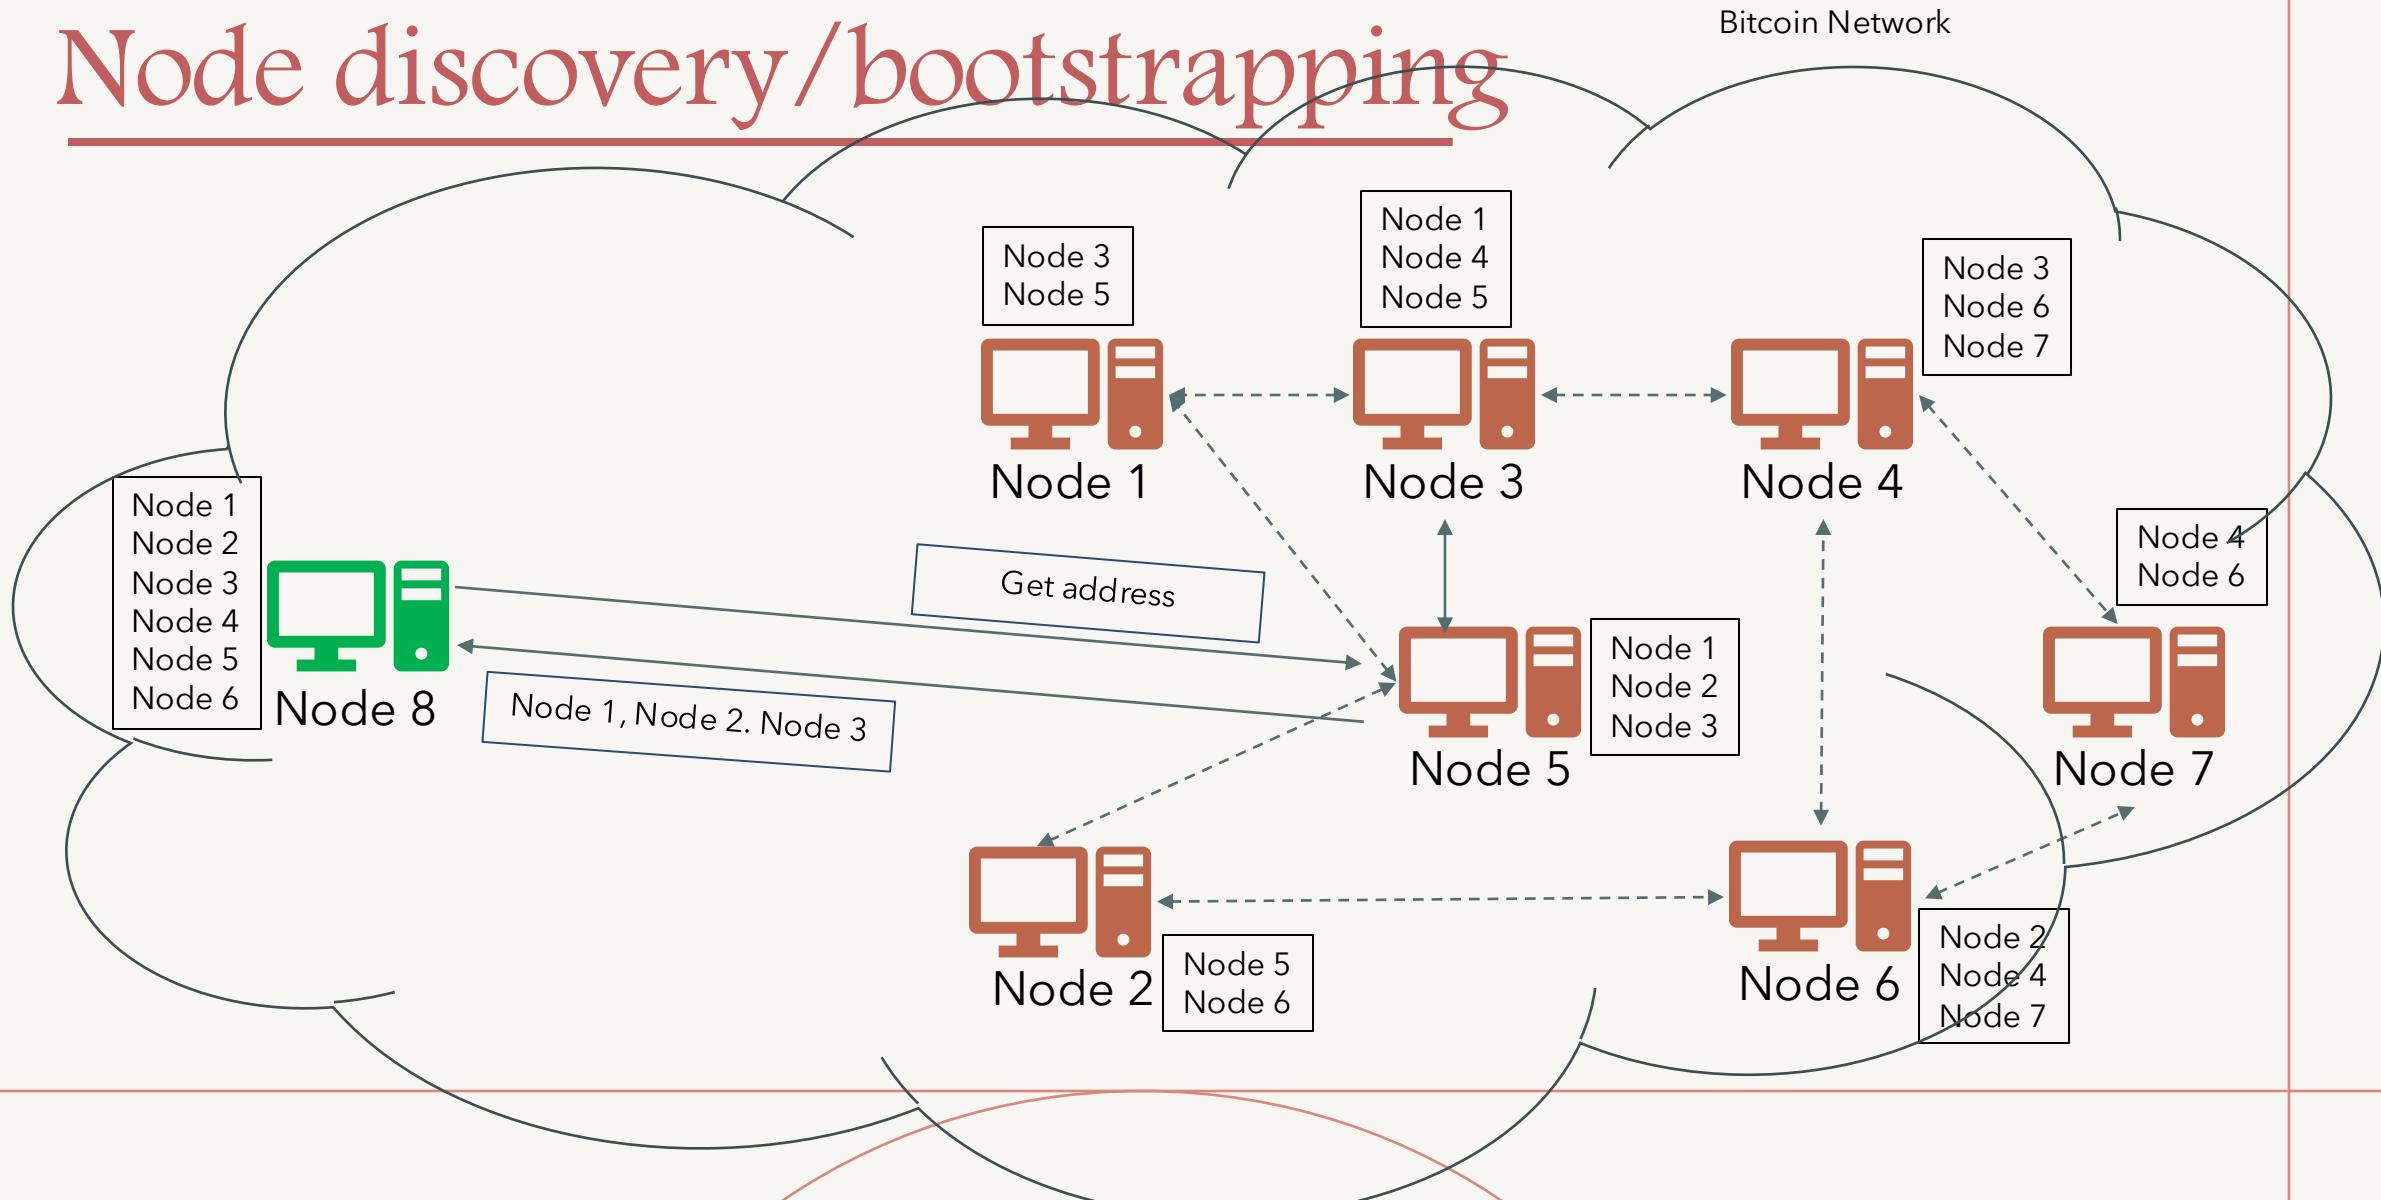
Node discovery/bootstrapping



Node discovery/bootstrapping



Node discovery/bootstrapping



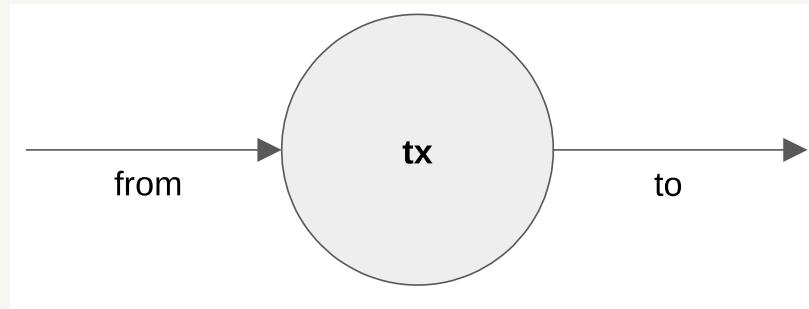
Bitcoin blockchain

- There are three different things to understand
 - Transaction
 - Block
 - Blockchain

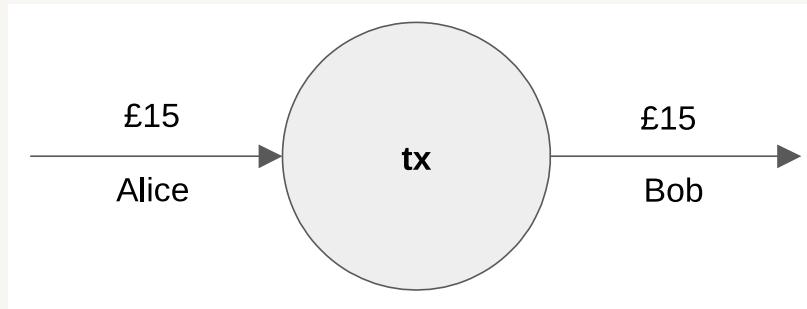
Transaction

- Transactions are the most important part of the bitcoin system
- Everything else in bitcoin is designed to ensure that transactions can be created, propagated on the network, validated, and finally added to the global ledger of transactions (the blockchain)
- Transactions are data structures that encode the transfer of values between participants in the bitcoin system
 - A transaction transfers money from someone to somebody else
- Each transaction is a public entry in bitcoin's blockchain, the global double-entry bookkeeping ledger
 - an entry that affects at least two different accounts
- All transactions are public
- Everybody can see all transactions in a blockchain explorer

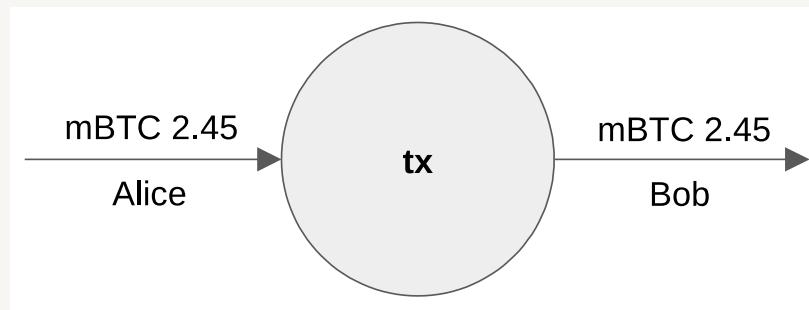
Transaction



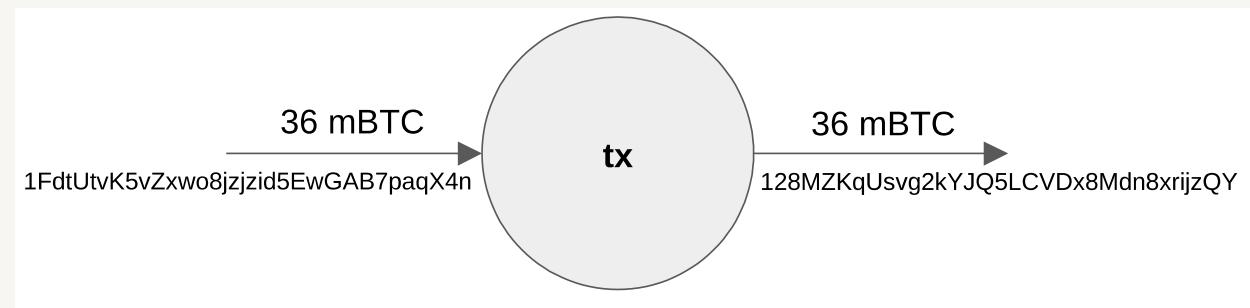
Abstract format of any transaction



A traditional transaction



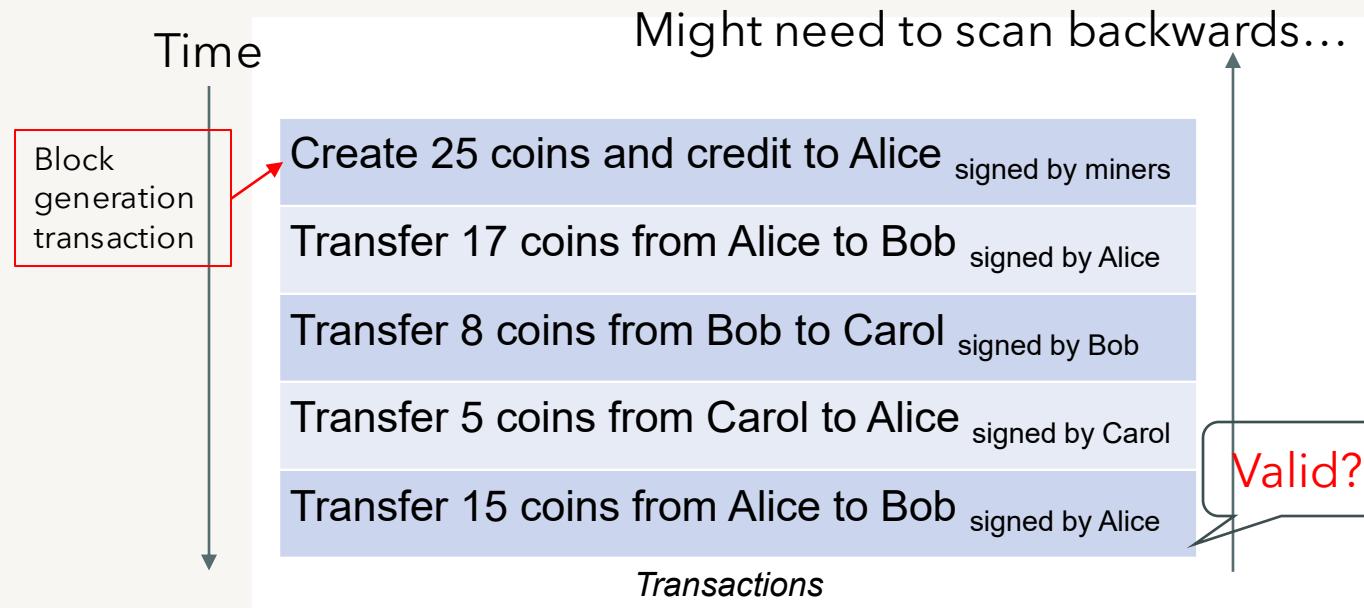
Abstract format of a bitcoin transaction



A bitcoin transaction

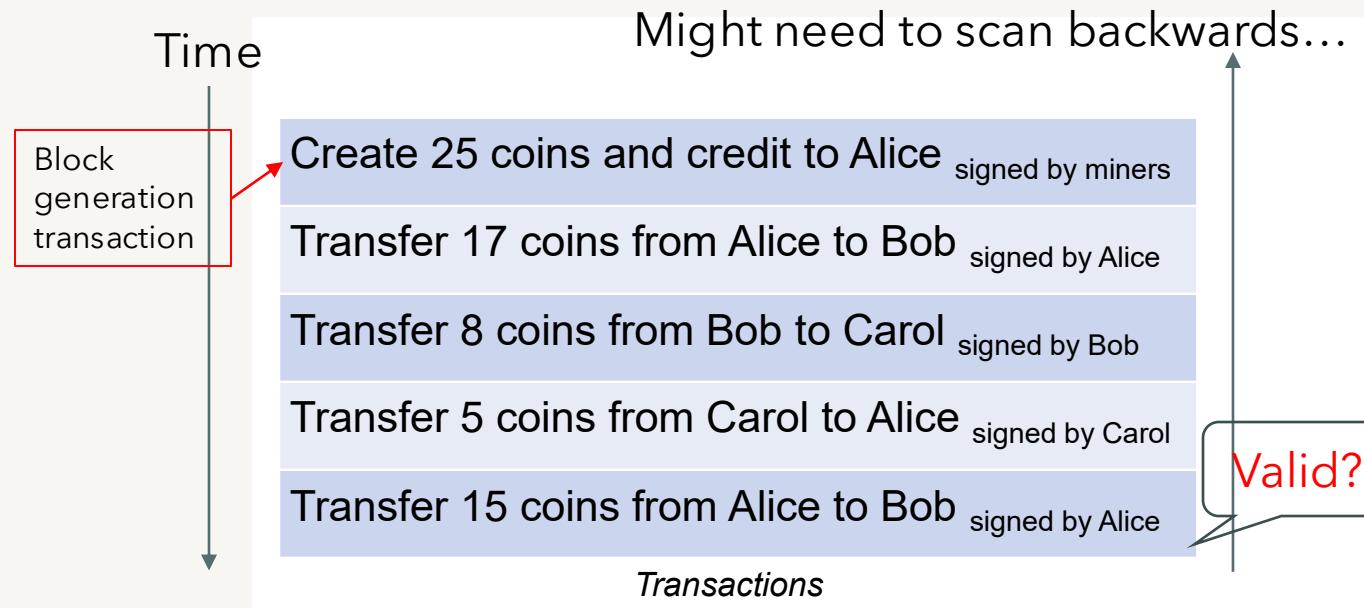


Account/transaction based Ledger?



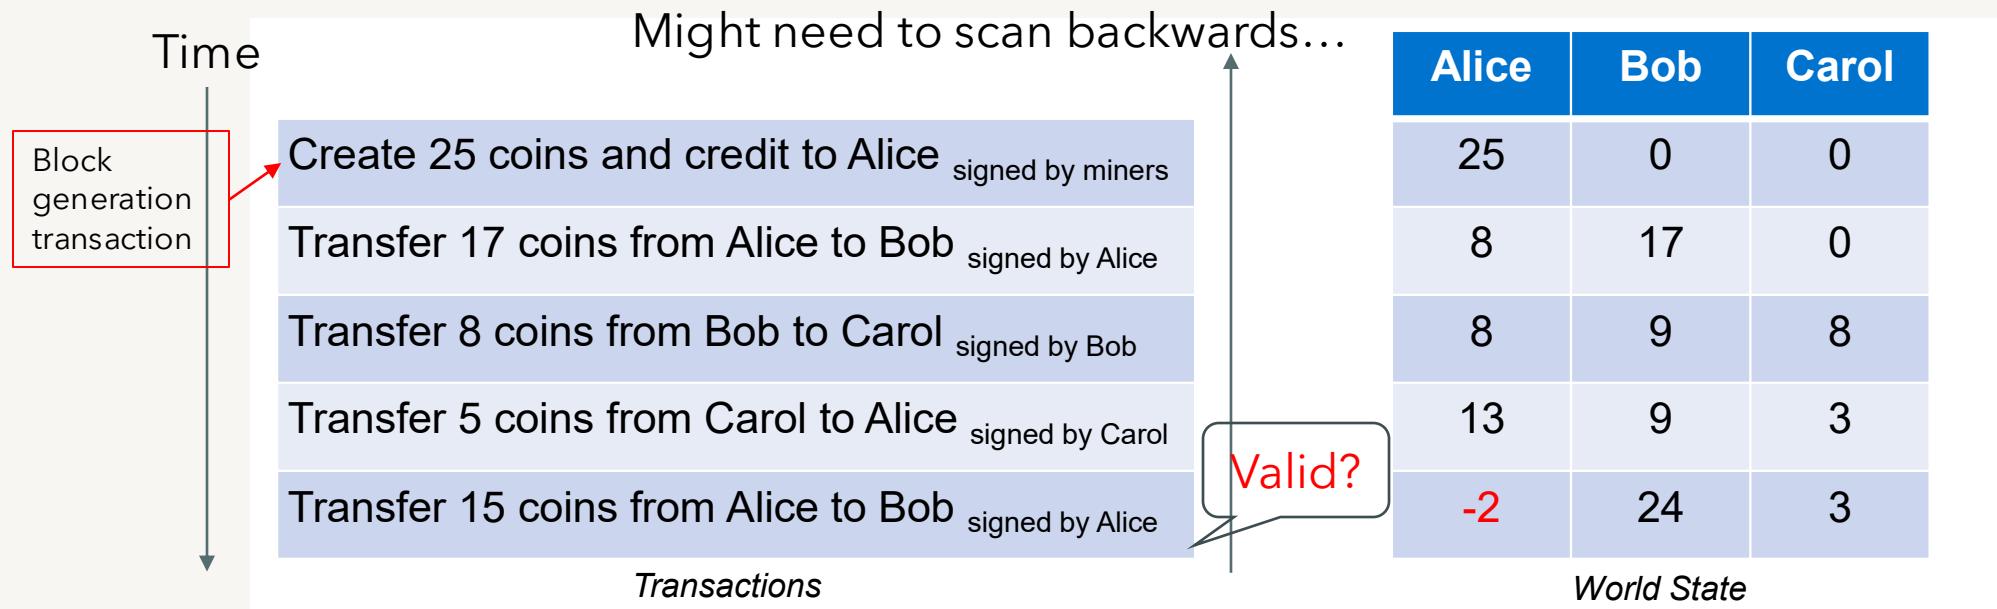
- *Intuitively:* At first, we consider Bitcoin to use an account-based ledger. However, an account-based approach takes a lot of effort to track the balances of every account
- In an account-based ledger, transactions can transfer arbitrary amounts of coins between accounts
- Transactions lead to a “world-state” of accounts and account balances

Account/transaction based Ledger?



- To validate a certain transaction, you might need to track very old transactions
- Since you do not know which old transaction, you need track each of the previous transactions one by one until you find the desired ones

Account/transaction based Ledger?

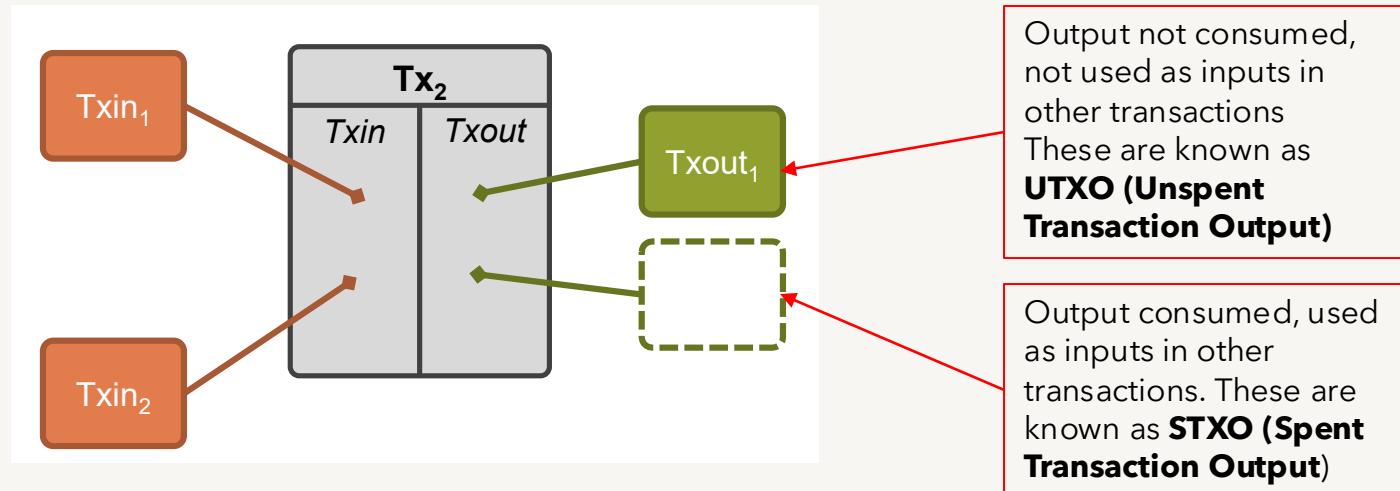


- One option is to maintain the world state in a separate database
 - Remember, in the world state you store the account balance
- That would require to maintain two separate databases: world state and transactions

Account/transaction based Ledger?

- Bitcoin's solution: a transaction-based ledger
- By using a transaction-based ledger, Bitcoin enables wallet owners to define conditional transactions using Bitcoin Script

Transaction Based Ledger



- Transactions (Tx) have a number of inputs and a number of outputs
 - Inputs (Txin): Former outputs, that are being consumed
 - Outputs (Txout): New outputs transferring the value
- In transactions (coinbase transaction) where new coins are created, no Txin is used (no coins are consumed)
- Each transaction has a unique identifier (TxID). Each output has a unique identifier within a transaction
- We refer to them (in this example) as #TX[#txout], e.g., 1[1], which is the second Txout of the second transaction

Transaction Based Ledger

This is known as a coinbase transaction

Create 25 coins and credit to Alice signed by miners

Transfer 17 coins from Alice to Bob signed by Alice

Transfer 8 coins from Bob to Carol signed by Bob

Transfer 5 coins from Carol to Alice signed by Carol

Transfer 15 coins from Alice to Bob signed by Alice

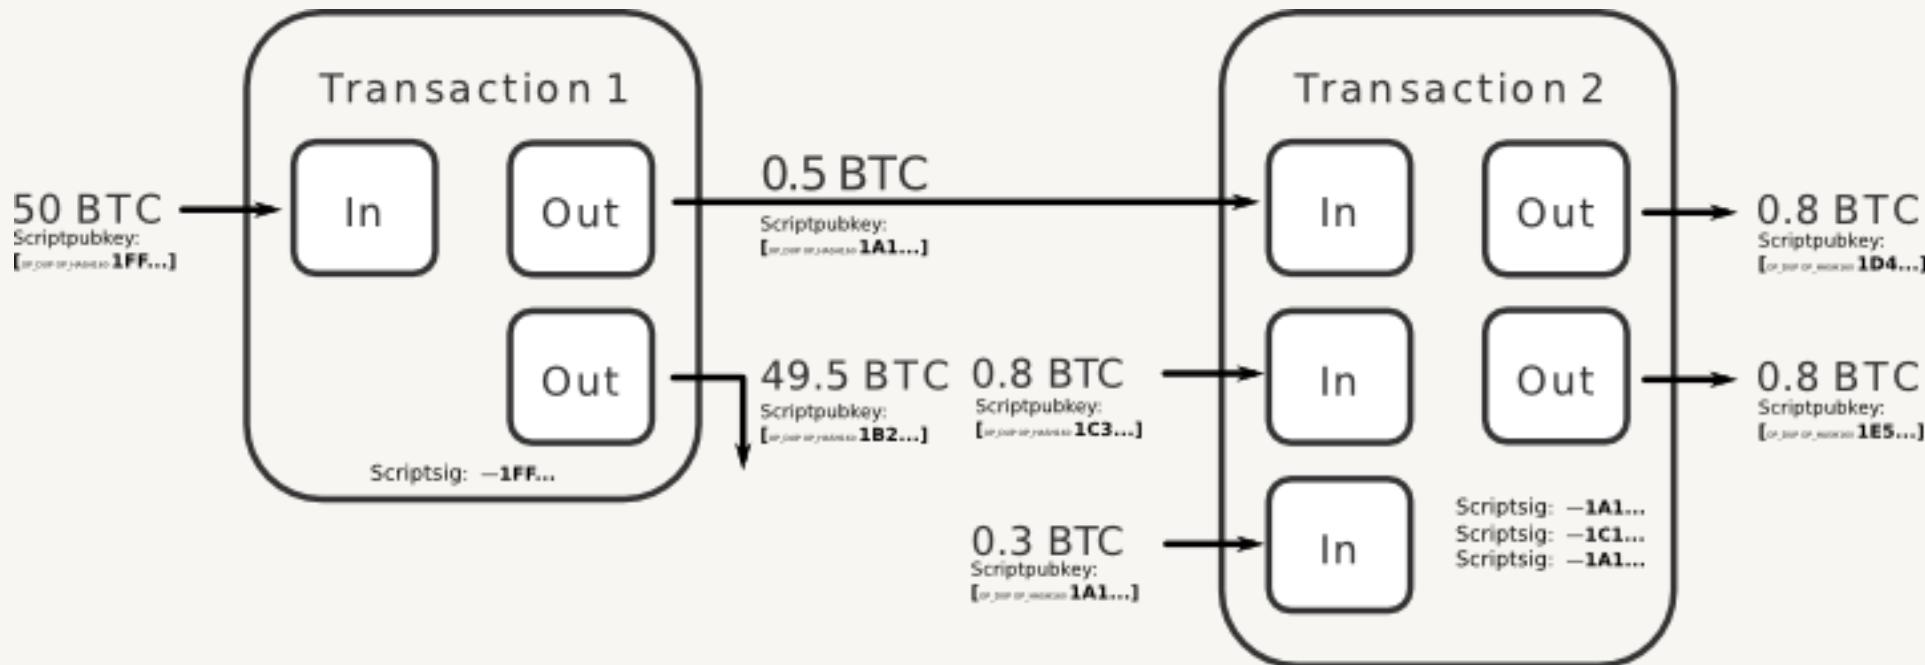
Transactions

| | |
|---|---|
| 0 | Txin: \emptyset Txout: 25.0 \rightarrow Alice <small>signed by the miner/system</small> |
| 1 | Txin: 0[0] Txout: 17.0 \rightarrow Bob, 8.0 \rightarrow Alice <small>signed by Alice</small> |
| 2 | Txin: 1[0] Txout: Txout: 8.0 \rightarrow Carol, 9.0 \rightarrow Bob <small>signed by Bob</small> |
| 3 | Txin: 2[0] Txout: 5.0 \rightarrow Alice, 3.0 \rightarrow Carol <small>signed by Carol</small> |
| 4 | Txin: 1[1], 3[0] Txout: 15.0 \rightarrow Bob, ? \rightarrow Alice <small>signed by Alice</small> |

Transaction Based Ledger

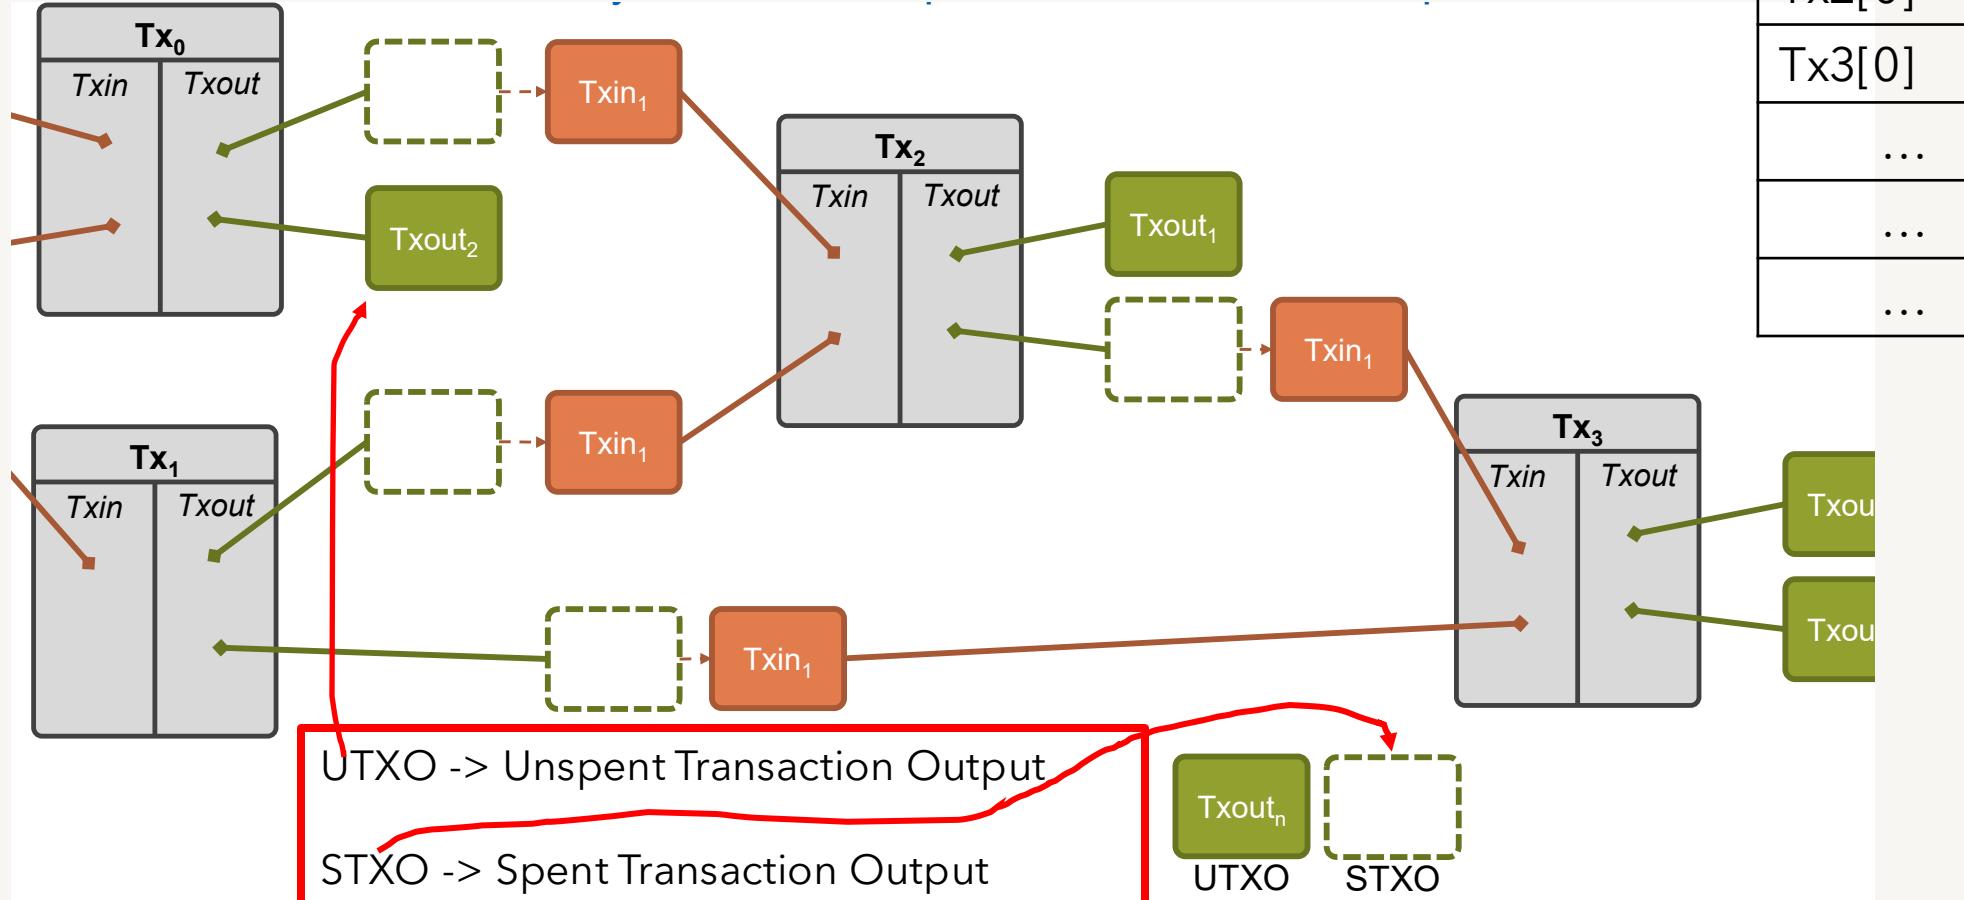
| | |
|--------------------------------------|--|
| Create 25 coins and credit to Alice | signed by miners |
| Transfer 17 coins from Alice to Bob | signed by Alice |
| Transfer 8 coins from Bob to Carol | signed by Bob |
| Transfer 5 coins from Carol to Alice | signed by Carol |
| Transfer 15 coins from Alice to Bob | signed by Alice |
| <i>Transactions</i> | |
| Joined payment | |
| 0 | Txin: \emptyset Txout: 25.0 \rightarrow Alice |
| 1 | Txin: 0[0] Txout: 17.0 \rightarrow Bob, 8.0 \rightarrow Alice |
| 2 | Txin: 1[0] Txout: Txout: 8.0 \rightarrow Carol, 9.0 \rightarrow Bob |
| 3 | Txin: 2[0] Txout: 5.0 \rightarrow Alice, 3.0 \rightarrow Carol |
| 4 | Txin: 1[1], 2[0] Txout: 15.0 \rightarrow Bob, ? \rightarrow Alice |

Connecting different transactions



All UTXOs are maintained
in an UTXO Table

Connecting different transactions



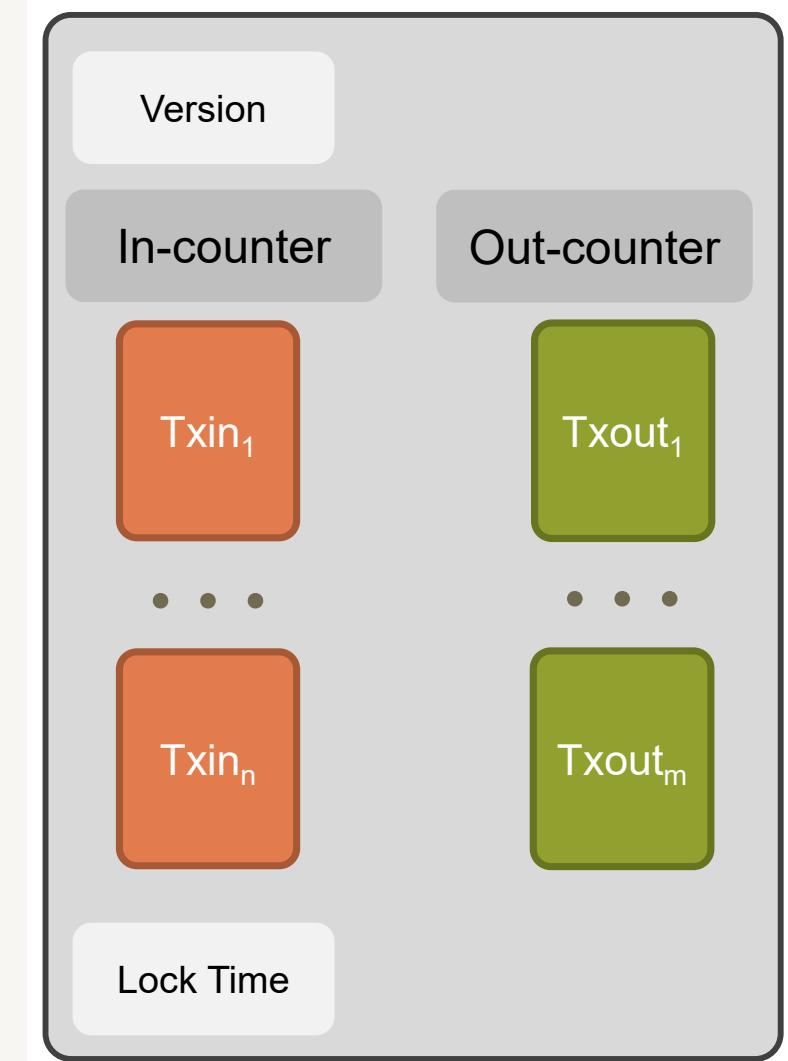
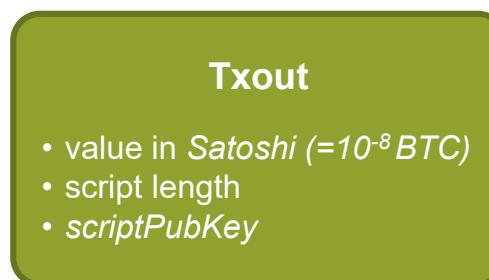
Transaction

- Each transaction has a list of inputs and outputs
- All inputs reference an existing unspent output or a coinbase transaction
- Inputs and outputs contain scripts (`scriptSig`, `scriptPubKey`) for verification and other metadata
- `lock_time`: is the time at which a particular transaction can be added to the blockchain, 0 means now

Input format



Output format



Transaction

- All inputs reference an existing unspent output or a

General format of a Bitcoin transaction (inside a block)

Version

In-counter

Out-counter

<https://en.bitcoin.it/wiki/Transaction>

| Field | Description | Size |
|-----------------|--|----------------------------|
| Version no | currently 1 | 4 bytes |
| In-counter | positive integer VI = VarInt | 1 - 9 bytes |
| list of inputs | <u>the first input of the first transaction is also called "coinbase" (its content was ignored in earlier versions).</u> | <in-counter>-many inputs |
| Out-counter | positive integer VI = VarInt | 1 - 9 bytes |
| list of outputs | the outputs of the first transaction spend the mined bitcoins for the block | <out-counter>-many outputs |
| lock_time | if non-zero and sequence numbers are < 0xFFFFFFFF: block height or timestamp when transaction is final | 4 bytes |

- previous transaction hash
- previous Txout-index
- script length
- scriptSig*

- value in *Satoshi* ($=10^{-8}$ BTC)
- script length
- scriptPubKey*

Lock Time

Transaction

It contains the size of the transaction, the number of inputs and outputs, the version and a lock-time. The hash is the transaction ID (TxID) discussed later

An array of all inputs. Each input contains the previous transaction hash (TxID) and the index of Txout. Also a signature script (*scriptSig*) is provided.

An array of all outputs. One output has two fields: the amount of the transferred coins and the *scriptPubKey*.



Transaction output

- An output contains instructions for sending bitcoins
 - Value is the number of Satoshi (1 BTC = 100,000,000 Satoshi) that this output will be worth when claimed
- There can be more than one output, and they share the combined value of the inputs
- If the input is worth 50 BTC but you only want to send 25 BTC,
 - Bitcoin will create two outputs worth 25 BTC: one to the receiver, and one back to you (known as "*change*", though you send it to yourself)
- Any input bitcoins not redeemed in an output is considered a *transaction fee*; whoever generates the block will get it

```
"vout": [
  {
    "value": 0.01500000,
    "scriptPubKey": "OP_DUP OP_HASH160 ab68025513c3dbd2f7b92a94e0581f5d50f654e7 OP_EQUALVERIFY OP_CHECKSIG"
  },
  {
    "value": 0.08450000,
    "scriptPubKey": "OP_DUP OP_HASH160 7f9b1a7fb68d60c536c2fd8aeaa53a8f3cc025a8 OP_EQUALVERIFY OP_CHECKSIG"
  }
]
```

<https://github.com/bitcoinbook/bitcoinbook/blob/develop/ch06.asciidoc>

Transaction output

- Almost all outputs are spendable chunks, known as UTXO (unspent transaction outputs)
 - The collection of all UTXO is known as the UTXO set or the UTXO table
- The UTXO set grows as new UTXO is created and shrinks when UTXO is consumed
- Every transaction represents a change in the UTXO set
- When we say that a user's wallet has "received" bitcoin
 - what we mean is that the wallet has detected an UTXO that can be spent with one of the keys controlled by that wallet

Transaction output

- The concept of a balance is created by the wallet application
 - The wallet calculates the user's balance by scanning the blockchain and aggregating the value of any UTXO the wallet can spend with the keys it controls
- Most wallets maintain a database or use a database service to store a quick reference set of all the UTXO they can spend with the keys they control
- Transaction outputs has another component:
 - A cryptographic puzzle that determines the conditions required to spend the output
 - The cryptographic puzzle is also known as a *locking script*, a *witness script*, or a *scriptPubKey*

